

Causal Discovery Algorithms

Agenda


- Discovery vs inference
- Causal discovery algorithms
 - PC
 - GES
 - LiNGAM
 - NOTEARS
 - Deep Learning
 - LLMs
- Code and examples

Discovery


vs

Inference


Graph



Directed graph




Directed graph




A directed graph (V, E) is:

- a set V (vertices, or nodes)
- and a set $E \subset V \times V$ (edges).


Directed graph




Directed graph



Directed acyclic graph (DAG)




Adjacency Matrix



1	0	1	0	0	0	0	0
2	0	0	1	0	0	0	1
3	0	0	0	1	1	1	1
4	0	0	0	0	0	0	0
5	0	0	0	0	0	1	1
6	0	0	0	0	0	0	1
7	0	0	0	0	0	0	0
	1	2	3	4	5	6	7

Topological order



Data Generation Process

$$X_1 = f_1(\varepsilon_1)$$

$$X_2 = f_2(X_1, \varepsilon_2)$$

$$X_3 = f_3(\varepsilon_3)$$

$$X_4 = f_4(X_2, X_3, \varepsilon_4)$$

$$X_5 = f_5(X_4, \varepsilon_5)$$

Structural Causal Model

$$X_1 = f_1(\varepsilon_1)$$


$$X_2 = f_2(X_1, \varepsilon_2)$$

$$X_3 = f_3(\varepsilon_3)$$

$$X_4 = f_4(X_2, X_3, \varepsilon_4)$$

$$X_5 = f_5(X_4, \varepsilon_5)$$

Causal Graph



Structural Causal Model

$$X_1 = f_1(\varepsilon_1)$$


$$X_2 = f_2(X_1, \varepsilon_2)$$

$$X_3 = f_3(\varepsilon_3)$$

$$X_4 = f_4(X_2, X_3, \varepsilon_4)$$

$$X_5 = f_5(X_4, \varepsilon_5)$$

Causal Discovery



Causal Inference

$$X_1 = f_1(\varepsilon_1)$$


$$X_2 = f_2(X_1, \varepsilon_2)$$

$$X_3 = f_3(\varepsilon_3)$$

$$X_4 = f_4(X_2, X_3, \varepsilon_4)$$

$$X_5 = f_5(X_4, \varepsilon_5)$$

Intervention



$$X_1 = f_1(\varepsilon_1)$$


$$X_2 = f_2(X_1, \varepsilon_2)$$

$$X_3 = f_3(\varepsilon_3)$$

$$X_4 = X_4$$

$$X_5 = f_5(X_4, \varepsilon_5)$$

Intervention



$$X_1 = f_1(\varepsilon_1)$$

$$X_2 = f_2(X_1, \varepsilon_2)$$

$$X_3 = f_3(\varepsilon_3)$$

$$X_4 = X_4$$

$$X_5 = f_5(X_4, \varepsilon_5)$$

Discovery vs Inference

- **Causal discovery:** finding the causal graph from data
- **Causal inference:** using the structural causal model to answer “what if” questions

Causality ladder

- **Association:** $E[Y|T=t]$
- **Intervention:** $E[Y|\text{do}(T=t)]$
- **Counterfactuals:** $E[Y|\text{do}(T=t), T=t']$




Causality ladder

- Association: $E[Y|T=t]$
- **Intervention**: $E[Y|\text{do}(T=t)]$
- Counterfactuals: $E[Y|\text{do}(T=t), T=t']$




Causality ladder

- Association: $E[Y|T=t]$
- Intervention: $E[Y|\text{do}(T=t)]$
- **Counterfactuals: $E[Y|\text{do}(T=t), T=t']$**



Causality ladder

- Association: $E[Y|T=t]$
- Intervention: $E[Y|\text{do}(T=t)]$
- **Counterfactuals: $E[Y|\text{do}(T=t), T=t']$**



Causality Ladder

- **Association**

If the patient received an aggressive treatment, it means his condition was already severe: the expected outcome is bad.

- **Intervention**

If we were to give all patients an aggressive treatment, the outcome would be good on average.

- **Counterfactual**

This specific patient received the non-aggressive treatment; this means his condition was mild; if we had given him the aggressive treatment, the outcome would have been good.

Simpson's paradox

		Condition		
		Mild	Severe	Total
Treatment	A	15% (210 / 1400)	30% (30 / 100)	16% (240 / 1500)
	B	10% (5 / 50)	20% (100 / 500)	19% (105 / 550)

Fundamental Problem

We often want to compute the average treatment effect,


$$ATE = E[Y|\text{do}(T=1)] - E[Y|\text{do}(T=0)],$$

but, for each subject, we either have $T=1$ or $T=0$, so we do not know the other.


Do Calculus

Do Calculus is a set of rules to compute, when possible, the effect interventions would have from observational data alone, given the causal graph.


Confounder



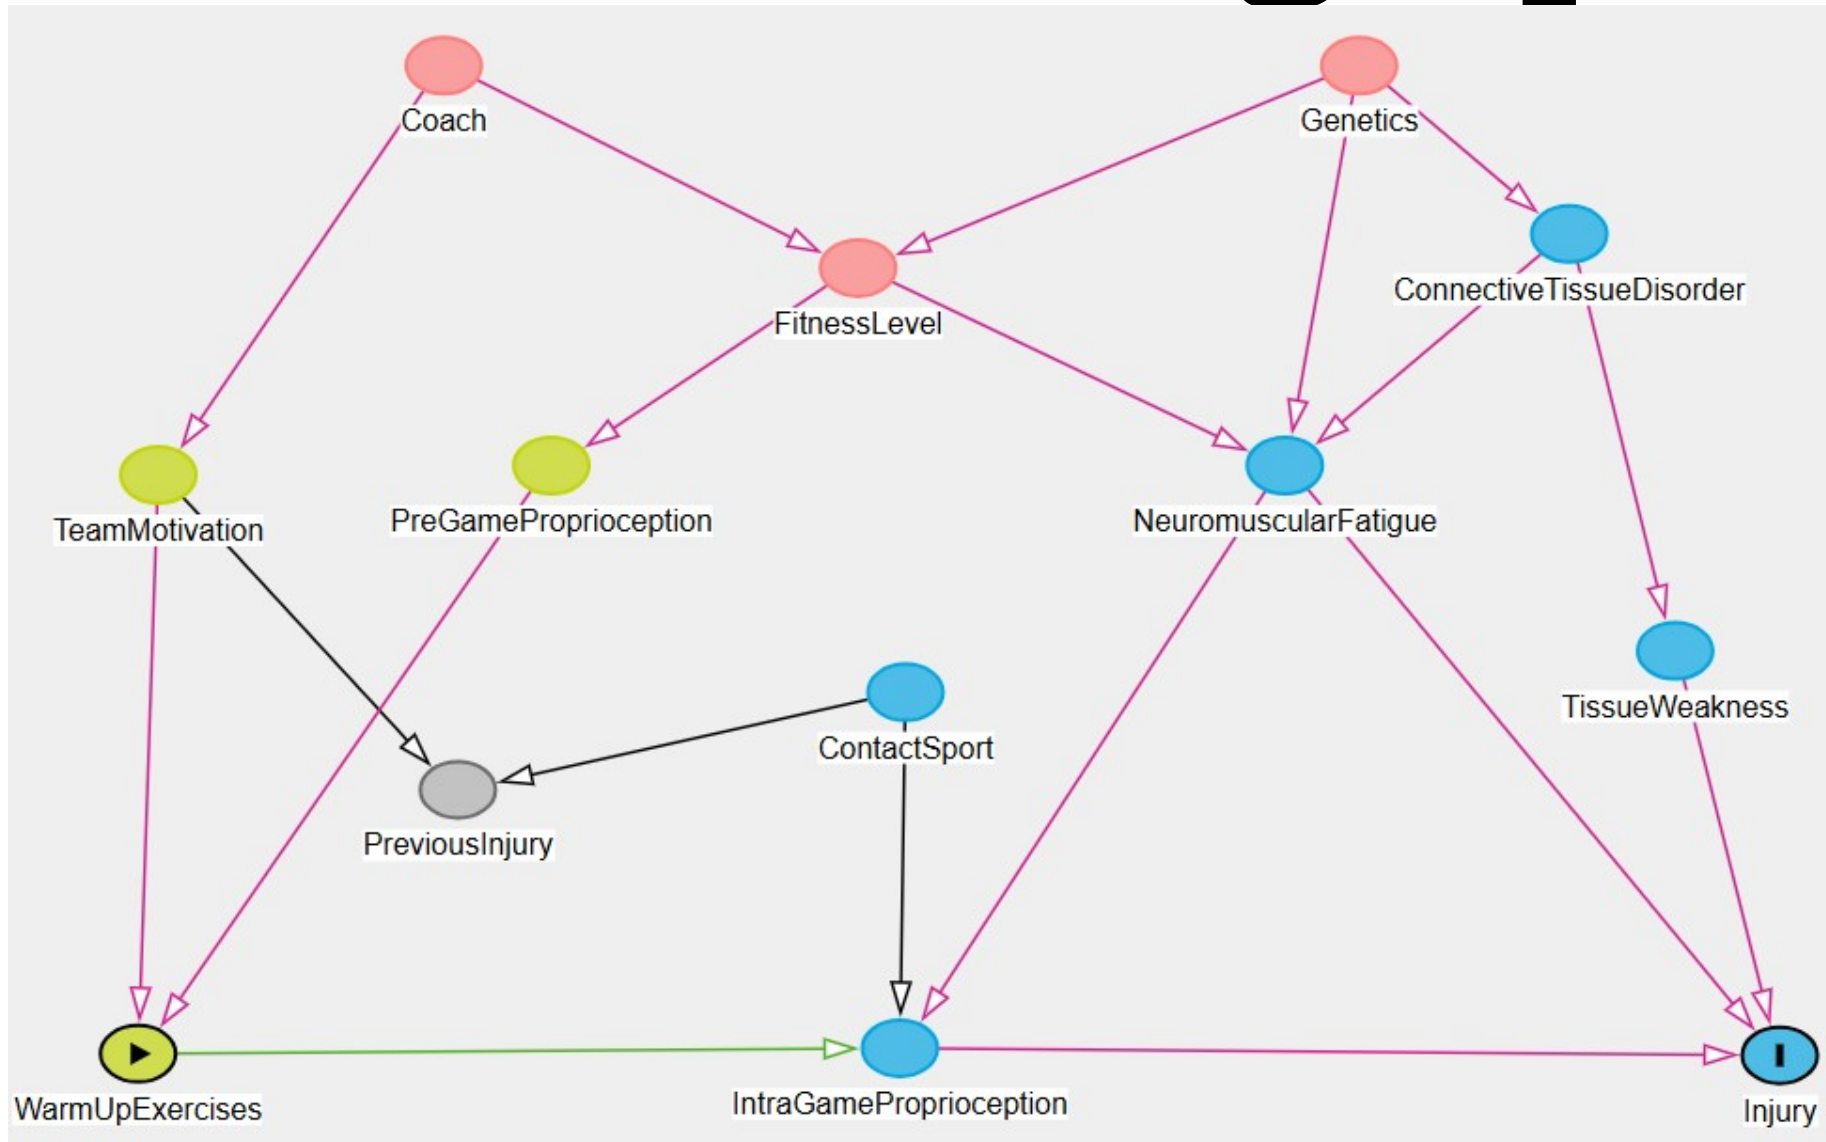
Collider



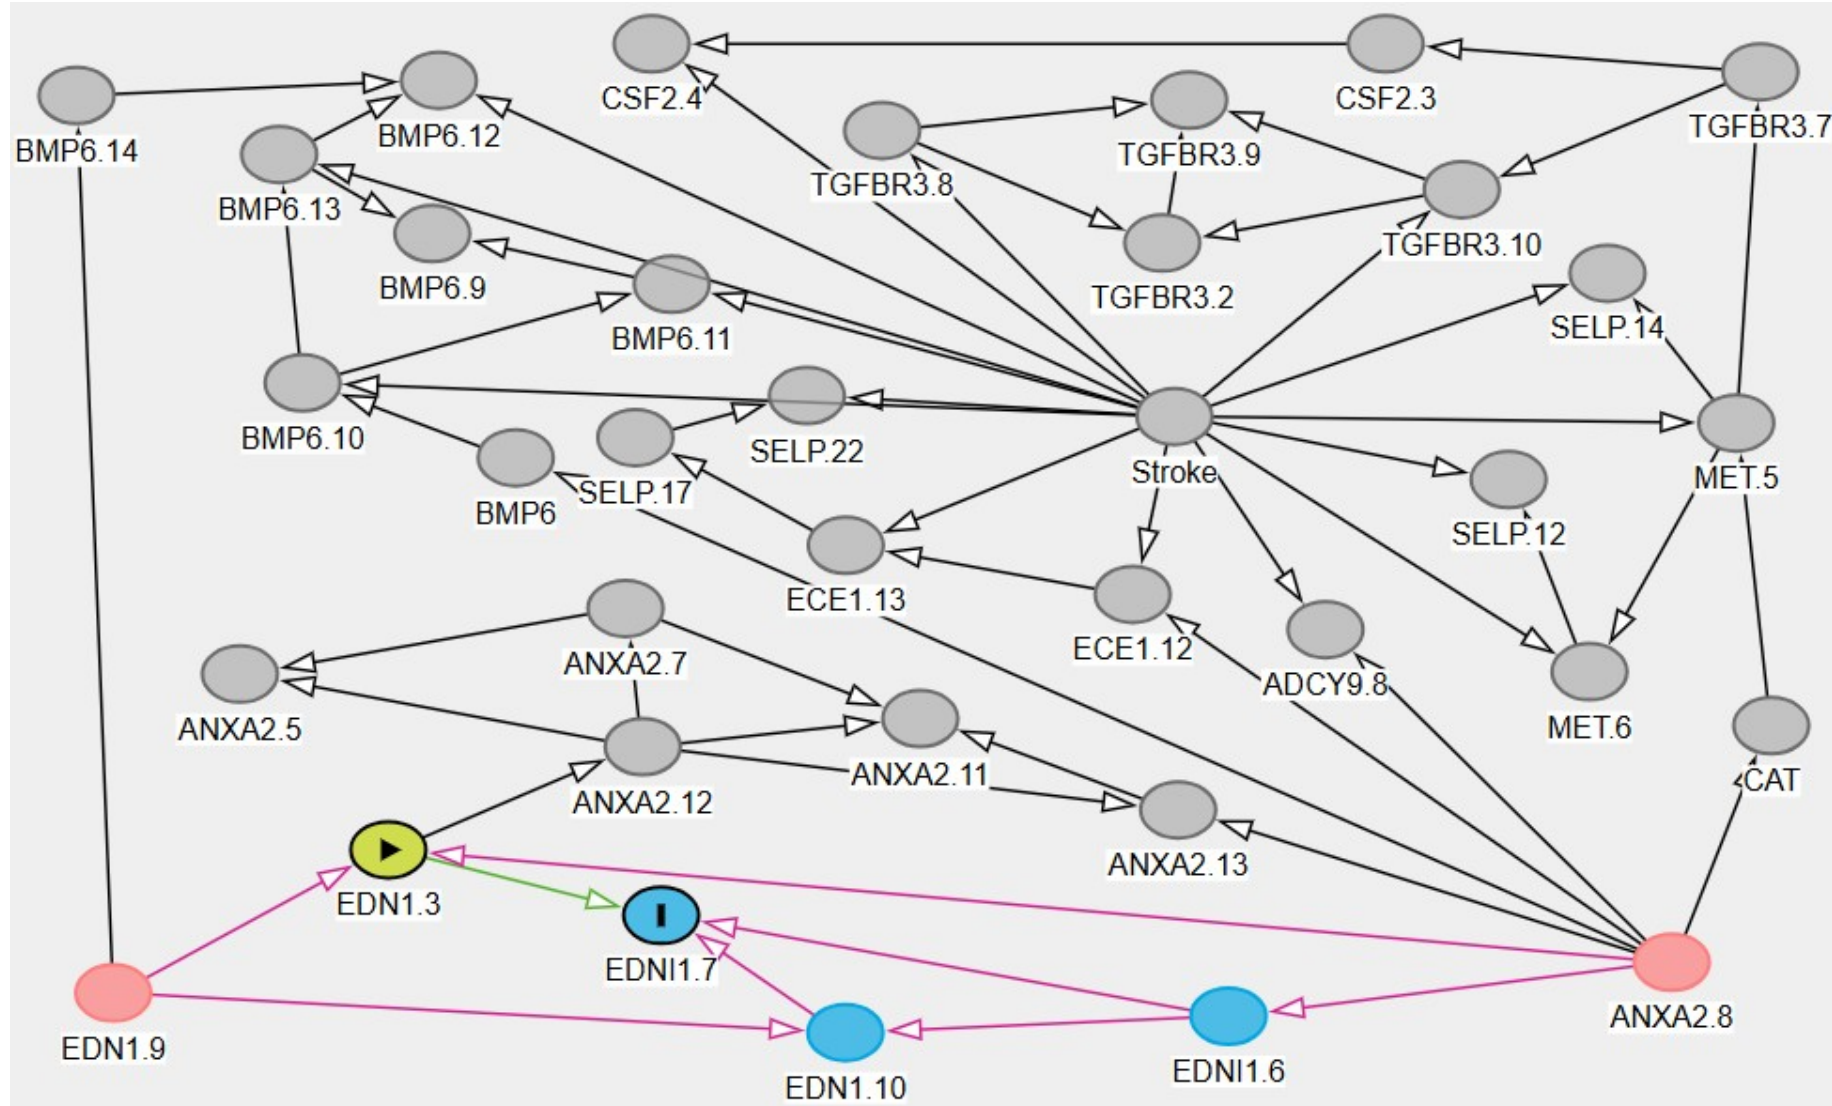
Mediator




Real-world causal graphs




Real-world causal graphs



Open and closed paths




open




open



open



closed




closed




closed



closed



open



open

○ Not in the conditioning set

● In the conditioning set

Causal inference

To assess the strength of the causal relation $X \rightarrow Y$:

- List all the (undirected) paths from X to Y
- All the non-causal paths should be blocked; if not, condition on one or more nodes to block them.
- All the causal paths should be open; if not, adjust the conditioning set to unblock them.

Causal Discovery Algorithms

Causal Discovery Algorithms


- **PC**: Conditional independence tests
- **GES**: Scores
- **LiNGAM**: Independent component analysis
- **NOTEARS**: Optimization
- Deep Learning
- LLMs

PC Algorithm


Conditional Independence



$$X \not\perp Y$$





$$X \perp Y$$
$$X \not\perp Y | Z$$




$$X \not\perp Y$$
$$X \perp Y | Z$$


- Start with a complete (undirected) graph
- Remove the edge $X-Y$ if $X \perp Y | Z$ for some (possibly empty) set of nodes Z
- For all $X-Z-Y$, if $X \perp Y$ and $X \not\perp Y | Z$, we have a collider $X \rightarrow Z \leftarrow Y$
- Propagate the orientation, assuming we have found all the colliders







- Start with a complete (undirected) graph
- Remove the edge $X-Y$ if $X \perp Y | Z$ for some (possibly empty) set of nodes Z
- For all $X-Z-Y$, if $X \perp Y$ and $X \not\perp Y | Z$, we have a collider $X \rightarrow Z \leftarrow Y$
- Propagate the orientation, assuming we have found all the colliders







- Start with a complete (undirected) graph
- Remove the edge $X-Y$ if $X \perp Y | Z$ for some (possibly empty) set of nodes Z
- For all $X-Z-Y$, if $X \perp Y$ and $X \not\perp Y | Z$, we have a collider $X \rightarrow Z \leftarrow Y$
- Propagate the orientation, assuming we have found all the colliders







- Start with a complete (undirected) graph
- Remove the edge $X-Y$ if $X \perp Y | Z$ for some (possibly empty) set of nodes Z
- For all $X-Z-Y$, if $X \perp Y$ and $X \not\perp Y | Z$, we have a collider $X \rightarrow Z \leftarrow Y$
- Propagate the orientation, assuming we have found all the colliders







- Start with a complete (undirected) graph
- Remove the edge $X-Y$ if $X \perp Y | Z$ for some (possibly empty) set of nodes Z
- For all $X-Z-Y$, if $X \perp Y$ and $X \not\perp Y | Z$, we have a collider $X \rightarrow Z \leftarrow Y$
- Propagate the orientation, assuming we have found all the colliders







- Start with a complete (undirected) graph
- Remove the edge $X-Y$ if $X \perp Y | Z$ for some (possibly empty) set of nodes Z
- For all $X-Z-Y$, if $X \perp Y$ and $X \not\perp Y | Z$, we have a collider $X \rightarrow Z \leftarrow Y$
- Propagate the orientation, assuming we have found all the colliders







- Start with a complete (undirected) graph
- Remove the edge $X-Y$ if $X \perp Y | Z$ for some (possibly empty) set of nodes Z
- For all $X-Z-Y$, if $X \perp Y$ and $X \not\perp Y | Z$, we have a collider $X \rightarrow Z \leftarrow Y$
- Propagate the orientation, assuming we have found all the colliders







- Start with a complete (undirected) graph
- Remove the edge $X-Y$ if $X \perp Y | Z$ for some (possibly empty) set of nodes Z
- For all $X-Z-Y$, if $X \perp Y$ and $X \not\perp Y | Z$, we have a collider $X \rightarrow Z \leftarrow Y$
- Propagate the orientation, assuming we have found all the colliders



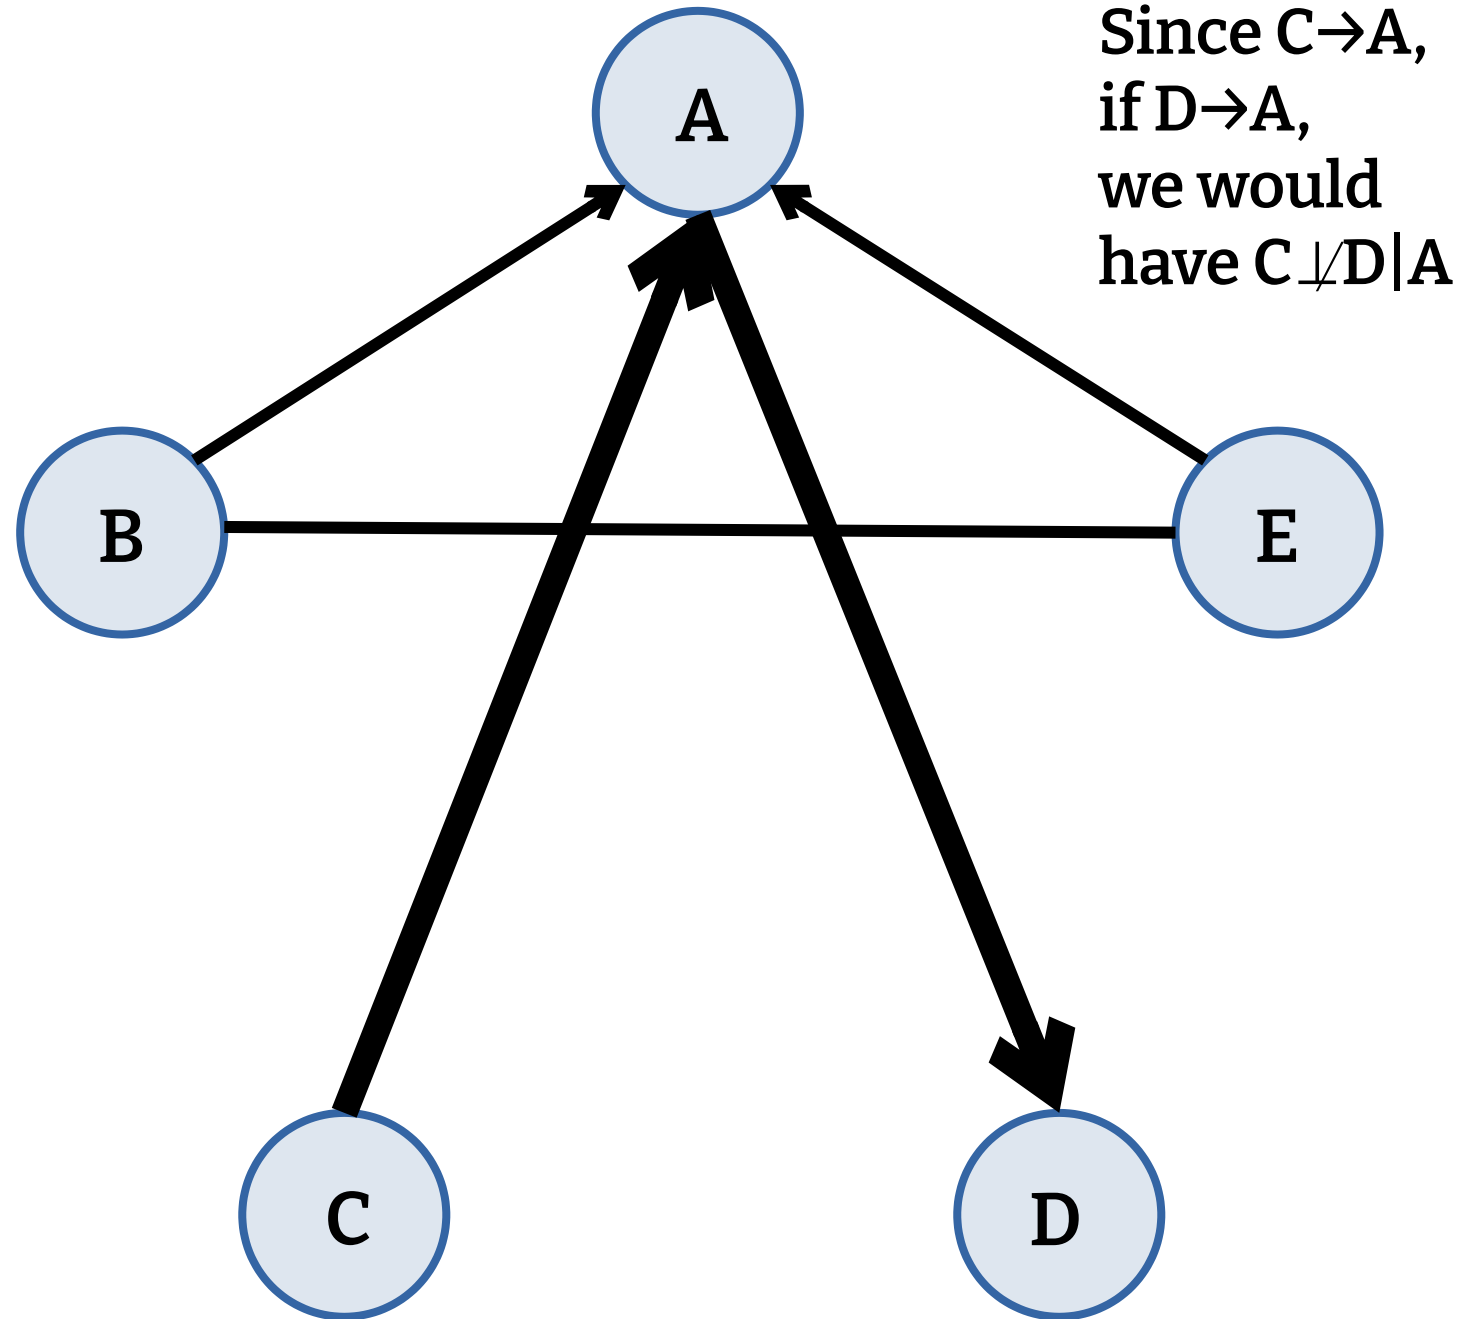



- Start with a complete (undirected) graph
- Remove the edge $X-Y$ if $X \perp Y | Z$ for some (possibly empty) set of nodes Z
- For all $X-Z-Y$, if $X \perp Y$ and $X \not\perp Y | Z$, we have a collider $X \rightarrow Z \leftarrow Y$
- Propagate the orientation, assuming we have found all the colliders






- Start with a complete (undirected) graph
- Remove the edge $X-Y$ if $X \perp Y | Z$ for some (possibly empty) set of nodes Z
- For all $X-Z-Y$, if $X \perp Y$ and $X \not\perp Y | Z$, we have a collider $X \rightarrow Z \leftarrow Y$
- Propagate the orientation, assuming we have found all the colliders







- Start with a complete (undirected) graph
- Remove the edge $X-Y$ if $X \perp Y | Z$ for some (possibly empty) set of nodes Z
- For all $X-Z-Y$, if $X \perp Y$ and $X \not\perp Y | Z$, we have a collider $X \rightarrow Z \leftarrow Y$
- Propagate the orientation, assuming we have found all the colliders




PC algorithm





Ground truth









Complete graph



Independence tests



Conditional independence tests



Colliders



Non-collider

Undirected edges

Final result

Markov equivalence class

Conditional independence relations cannot separate all DAGs: they can only recover the Markov equivalence class (MEC) of the causal model.



Independence tests

- Partial correlation (Fisher Z)
- Kernel-based tests
- χ^2 test (for discrete data)

GES

GES (Greedy Equivalent Search)

- **Start with an empty graph**
- **Greedily add the edges whose addition increases the score the most**
- **Greedily remove the edges whose removal increase the score the most**
- **Score: BIC, when forecasting a variable from its parents**

LiNGAM

LinGAM

The structural causal model (SCM)

$$X_1 = \mu_1 + \varepsilon_1$$

$$X_2 = \mu_2 + a_{21} X_1 + \varepsilon_2$$

$$\vdots$$

$$X_k = \mu_k + a_{k1} + \dots + a_{kk-1} X_{k-1} + \varepsilon_k$$

can be written $X = \mu + AX + \varepsilon$, or $\varepsilon = (I - A)X - \mu$, with $\forall i \neq j \ \varepsilon_i \perp \varepsilon_j$

If the ε_i 's are non-Gaussian, ICA (independent component analysis) can recover the linear transformation $(I - A)X$ by looking for the directions in which the data is the least Gaussian.

NOTEARS

NOTES

Find a matrix W such that $WX \approx X$ and the corresponding graph is acyclic.

The acyclicity condition can be written

$$\text{trace } \exp |W| = n$$

(where $|\cdot|$ is the elementwise absolute value)

$$\exp A = I + A + \frac{1}{2!} A^2 + \frac{1}{3!} A^3 + \dots$$

$\text{diag } A^k$: number of cycles of length k

NOTEARS


Find a matrix W such that $WX \approx X$ and the corresponding graph is acyclic.

$$\begin{array}{ll} \text{Find} & A \\ \text{To minimize} & \text{Mean}_i \|X_i - AX_i\|_F^2 + \lambda \|A\|_1 \\ \text{Such that} & \text{Trace } e^{|A|} = d \end{array}$$

Variants of those algorithms


- **Unobserved confounders**
- **Interventions**
- **Cycles**
- **Time series**

Unobserved Confounders




$$X_1 \perp X_4 \mid X_3$$

Unobserved Confounders




$$X_1 \perp X_4 \mid X_3$$




$$X_1 \not\perp X_4 \mid X_3$$

Unobserved Confounders



$$X_1 \perp X_4 \mid X_3$$

Instrumental variables



$$X_1 \not\perp X_4 \mid X_3$$

Deep Learning

Deep Learning

- Replace linear transformations with neural networks
- Better search algorithms for score-based methods (DP, A^*)
- First look for a topological order
- Reinforcement learning for search
- Reinforcement learning to progressively build the solution

Beyond NOTEARS

- NOTEARS is linear:

Find A

To minimize $\text{Mean}_i \|X_i - AX_i\|_F^2 + \lambda \|A\|_1$

Such that $\text{Trace } e^{A \odot A} = d$

Beyond NOTEARS

- It can be made non-linear:

Find	A
	g_1, g_2 (pointwise)
To minimize	$\text{Mean}_i \left\ X_i - g_2 \left(A g_1 (X_i) \right) \right\ _F^2 + \lambda \ A\ _1$
Such that	$\text{Trace } e^{A \odot A} = d$

Beyond NOTEARS

- It can be made non-linear:

Find g_W neural net

To minimize $\text{Mean}_i \|X - g_W(X)\|^2 + \lambda \|W\|^2$

Such that $\text{Trace } e^C = d$

Where $C = |W^{(L)}| \cdot |W^{(L-1)}| \dots |W^{(1)}|$
(neural network connectivity matrix)

Beyond NOTEARS

- Learn a binary mask (Gumbel softmax)

Find $A \in \{0, 1\}^{n \times n}$
 $g_i : \mathbf{R}^n \rightarrow \mathbf{R}$

Such that $X_i \approx g_i(A_{\cdot i} \odot X)$
 A acyclic

Beyond NOTEARS

- The model $X = AX + Z$ defines an auto-encoder

$$X = (I - A)^{-1} Z \quad \text{decoder}$$

$$Z = (I - A)X \quad \text{encoder}$$

- It can be made nonlinear

$$X = f_2[(I - A)^{-1} f_1(Z)] \quad \text{decoder}$$

$$Z = f_4[(I - A) f_3(X)] \quad \text{encoder}$$

- And trained as a VAE, with a NOTEARS-like penalty


Beyond GES

$$\text{Score} = \sum_{\substack{j \in \llbracket 1, d \rrbracket \\ \text{variable}}} \sum_{\substack{k \in \llbracket 1, n \rrbracket \\ \text{observation}}} \log p(x_{jk} | \text{Pa}_{jk}; \theta_j) - \frac{|\theta_j|}{2} \log n$$

Beyond GES

- GES is greedy: it is not an exact search
- Exhaustive search is not reasonable: there are too many DAGs
- It is a shortest path problem on the subset lattice: it can be solved with dynamic programming
- The lasso gives a consistent A^* heuristic

Beyond GES



$$2^n \ll n! \ll \#\text{DAGs}$$

Order Search

- First find a topological order, then the DAG
- CAM:
 - Neighbourhood selection (GAM Boosting)
 - Complete DAG (unpenalized GES)
 - Pruning (p-values of GAM terms)

Reinforcement Learning

- Stochastic search for a high-score DAG:
 - State: DAG
 - Action: new (nearby) DAG
 - Reward: score

Reinforcement Learning

- Stochastic search for a high-score DAG:
 - State: Bootstrap sample
 - Action: DAG
 - Reward: $\text{score} + \lambda \cdot 1_{\text{DAG}} + \mu \cdot \text{NOTEARS}$

Reinforcement Learning

- Progressively build a topological order:
 - State: embedding of the latest variable selected
 - Action: variable to add
 - Reward: BIC improvement
 - Encoder: self-attention $R_{n \times d} \rightarrow R_{n \times d}$
 - Decoder: LSTM
 - Optimization: actor critic, policy gradient

LLM

LLM

- Ask an LLM to give the causal graph, from **metadata alone** (column names + descriptions)
- Ask an LLM to scour the literature to extract causal relations, and put them in a database (knowledge graph).

LLM

- Ask an LLM to give the causal graph, from **metadata** alone (column names + descriptions)
- Ask an LLM to scour the literature to extract causal relations, and put them in a database (knowledge graph).

Code

gCastle

```
import castle.algorithms  
import networkx as nx
```

```
model = castle.algorithms.PC()  
model.learn(X)
```

```
A = model.causal_matrix  
A = pd.DataFrame( A, columns = X.columns, index = X.columns )  
g = nx.from_pandas_adjacency( A, create_using = nx.DiGraph )
```

causal-learn

```
from causallearn.search.ConstraintBased.PC import pc
from causallearn.utils.cit import fisherz, kci, chisq, gsq
```

```
# Computation
g = pc(X.values)
```

```
# Extract the adjacency matrix
A = pd.DataFrame( g.G.graph )
A = ( A == -1 ).astype(int)
A.columns = A.index = X.columns.copy()
```

causal-learn

```
from causallearn.search.ScoreBased.GES import ges
```

```
r = ges(X.values)
```

```
A = pd.DataFrame( r['G'].graph )
```

```
A = ( A == -1 ).astype(int)
```

```
A.columns = A.index = X.columns.copy()
```

causal-learn

```
from causallearn.search.FCMBased.lingam import ICALiNGAM
```

```
model = ICALiNGAM()
```

```
model.fit(X)
```

```
A = model.adjacency_matrix_
```

```
A = pd.DataFrame( A != 0, index = X.columns, columns = X.columns )
```

cdt

```
import cdt                                # Causal discovery toolbox
import networkx as nx

pc = cdt.causality.graph.PC()             # Continuous, Gaussian variables
g = pc.predict(d)

A = nx.adjacency_matrix(g).todense()
A = pd.DataFrame( X, index = g.nodes, columns = g.nodes )
```


dowhy

```
from dowhy import CausalModel
model = CausalModel(
    data = X,
    treatment = "T",
    outcome = "Y",
    graph = ' '.join( nx.generate_gml(g) ),
)
estimand = model.identify_effect()
estimate = model.estimate_effect(
    estimand,
    method_name = "backdoor.linear_regression",
)
model.refute_estimate(
    estimand, estimate,
    method_name = "random_common_cause",
)
```

Examples





Examples

- Different algorithms give very different results



Examples

- Different bootstrap samples give very different results



Conclusion

Summary

- **PC**: conditional independence test
- **GES**: goodness-of-fit of models predicting y from its parents
- **LiNGAM**: independent component analysis
- **NOTEARS**: optimization problem, with acyclicity constraint
- **Software**: gCastle, causal-learn, cdt

Conclusion

- Do not start from data, but from a domain knowledge causal graph
- Do not use a single causal discovery algorithm, but several
- Do not run them on just the data, but also on bootstrap samples
- Do not only look at the output, look at the ingredients of the algorithms

References

Introduction to Causal Inference (B. Neal, 2020)

A survey on causal discovery: theory and practice (A. Zanga and F. Stella, 2023)


<https://github.com/huawei-noah/trustworthyAI/tree/master/gcastle>

<https://causal-learn.readthedocs.io/en/latest/>

<https://cran.r-project.org/web/views/CausalInference.html#dag>

Extra Slides

Causality competition



ADIA : Lab

COMPETITION GRANTS ABOUT CRUNCH TESTIMONIALS

JOIN NOW

GRANTS

\$100,000

HOW **CAUSAL DISCOVERY** CAN HELP **RESOLVE CRITICAL CHALLENGES** IN AI RESEARCH?

Join the latest ADIA Lab x Crunch Foundation competition to build a groundbreaking AI framework. Get a chance to win up to \$100,000.

JOIN THE CRUNCH →




ADIA : Lab

**"TRUTH IS
COMPLETELY OUT
OF CONTROL IF
YOU DON'T
UNDERSTAND
WHY"**



ADIA : Lab

**"WHY IS
WHAT
SEPARATE
YOU FROM
THEM"**




ADIA : Lab

**"LUCKY
IS WHO HAS
BEEN ABLE TO
UNDERSTAND
THE CAUSE OF
THINGS"**

Backdoor adjustment

$$E[Y|\text{do}(T = t)] = E_W E[Y|T = t, W]$$

$$P[Y|\text{do}(T = t)] = \sum_w P[Y|t, w] P[w]$$



With linear models: fit a regression

$$Y = \alpha + \beta T + \gamma W$$

and average W out:

$$Y = \alpha + \beta T + \gamma E[W]$$

Kernel methods

- Many machine learning algorithms do not really require coordinates, but just the Gram matrix $K_{ij} = \langle \mathbf{x}_i, \mathbf{x}_j \rangle$.
- Increasing the dimension, $K_{ij} = \langle \phi(\mathbf{x}_i), \phi(\mathbf{x}_i) \rangle$ does not change the size of the Gram matrix.
- We do not even need to compute $\phi(\mathbf{x}_i)$: we just need a (positive definite) kernel, $\kappa(\mathbf{x}, \mathbf{y}) = \langle \phi(\mathbf{x}), \phi(\mathbf{y}) \rangle$.

BIC Score

$$\text{Score} = \sum_{\substack{j \in \llbracket 1, d \rrbracket \\ \text{variable}}} \sum_{\substack{k \in \llbracket 1, n \rrbracket \\ \text{observation}}} \log p(x_{jk} | \text{Pa}_{jk}; \theta_j) - \frac{|\theta_j|}{2} \log n$$

Local BIC score of $X \rightarrow y$

$$H = \log(\Sigma[y,y] - \Sigma[y,X] \Sigma[X,X]^{-1} \Sigma[X,y])$$

$$-BIC = n \cdot H + \lambda \cdot k \cdot \log(n)$$

n = number of observations

k = number of variables in X

Σ = variance matrix of the data


Markov equivalence class

A **CPDAG** (complete partially directed acyclic graph) is a PDAG where

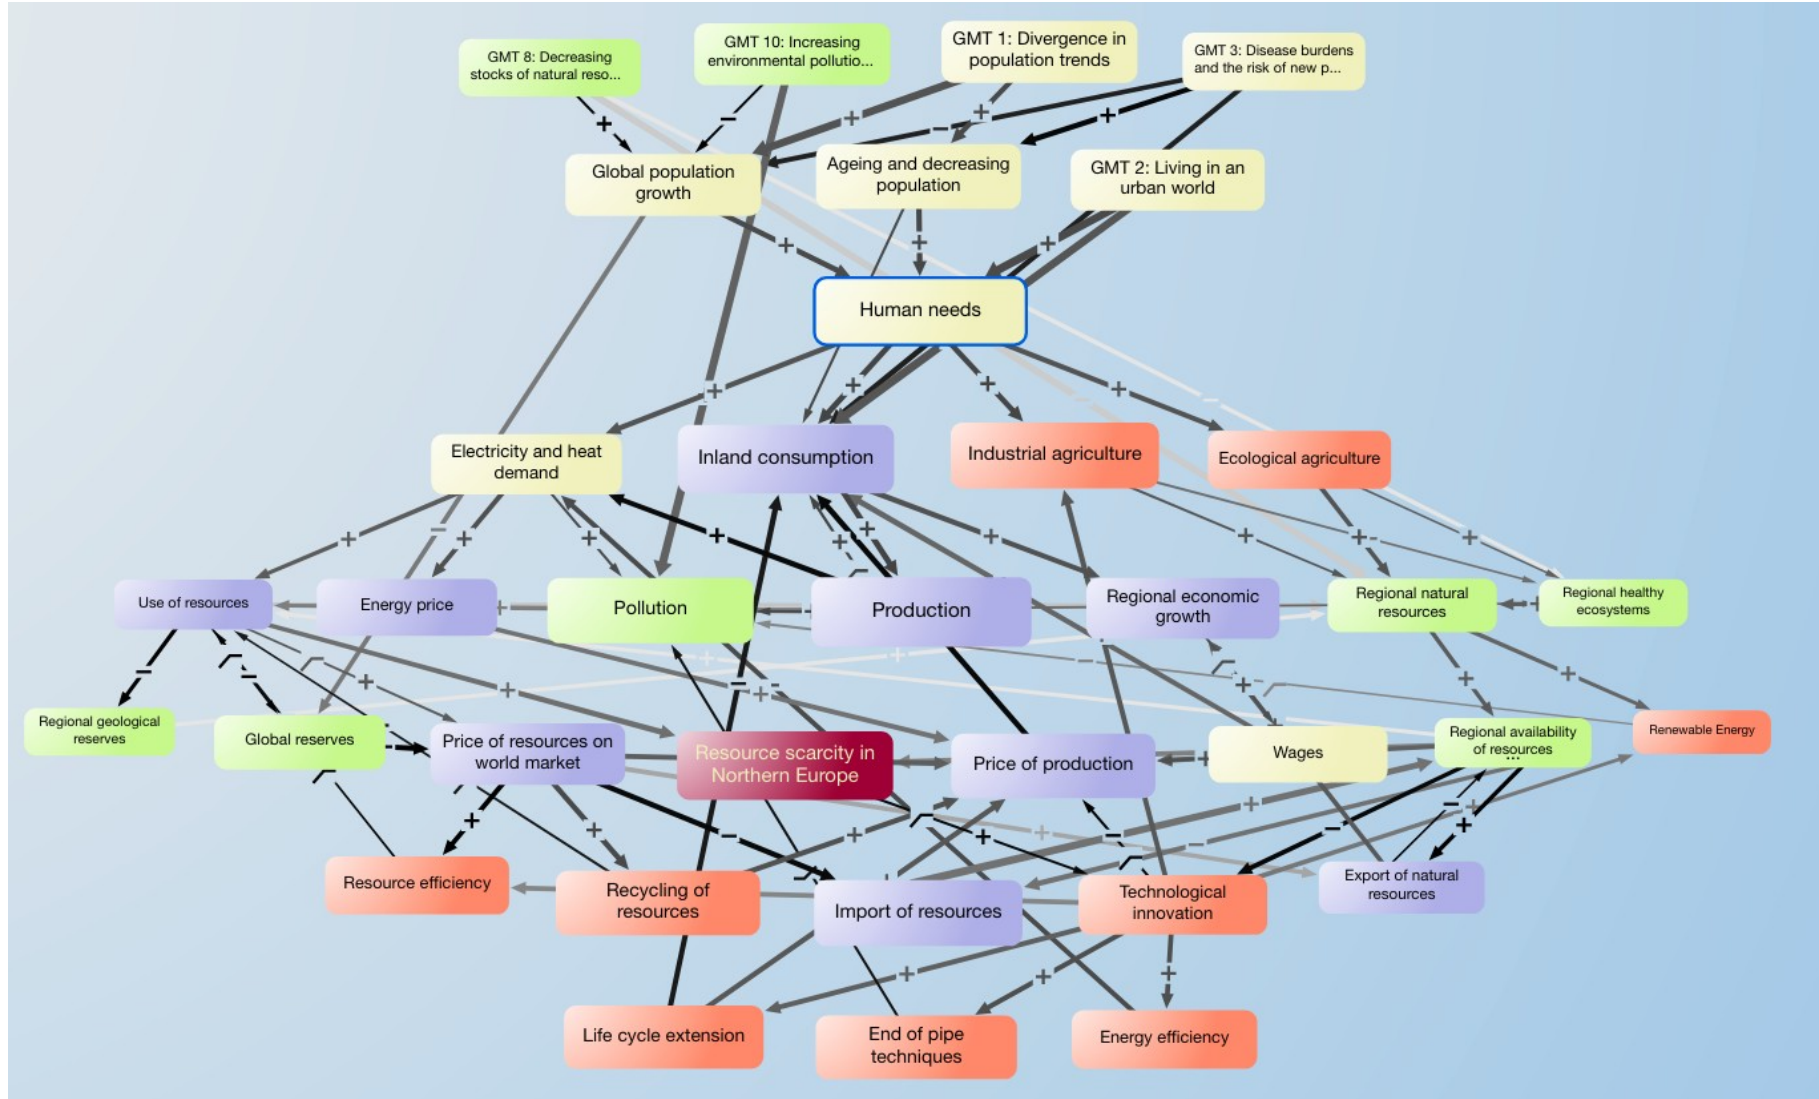
- All undirected edges are **reversible** (you can choose either direction, that does not change the MEC)
- All directed edges are **compelled** (if you flip them, the graph moves to another MEC)

Unused Slides

Real-world causal graphs

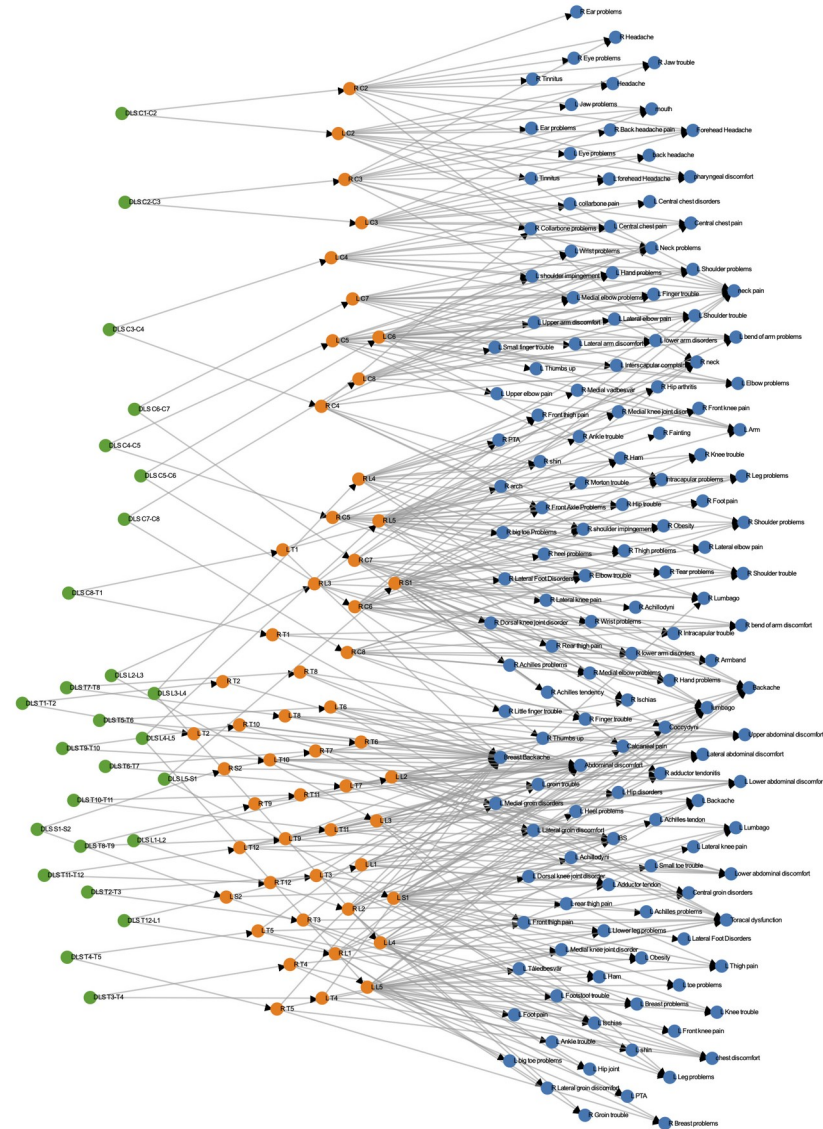


Real-world causal graphs



Impact assessment of global megatrends (U. Lorenz and H.V. Haraldson, 2014)

Real-world causal graphs




<https://arxiv.org/abs/1906.01732>

Variants of those algorithms


- **FCI: PC variant, allowing for unobserved confounders**
- **FGES: faster implementation of GES**
- **ARGES: another GES variant, for high-dimensional data**
- **GFCI: GES variant allowing for unobserved confounders (FCI on the FGES skeleton)**
- **CCD: PC/FCI with feedback (cycles)**
- **LiNG: LiNGAM with feedback**
- **Other LiNGAM variants: non-linear, post-non-linear**
- **CD-NOD**

Back-door adjustment



$$P[Y|\text{do}(T = t)] = \sum_w P[Y, t, w]P[w]$$

Front-door adjustment




$$P[Y|\text{do}(t)] = \sum_m P[m|t] \sum_{t'} P[y|m, t'] P[t']$$


Causal inference

- The sufficient conditioning set is not unique.
- The parents of X form a sufficient conditioning set, but it may be needlessly large.
- More generally, a sufficient conditioning set is a set of nodes blocking all the non-causal paths from X to Y , and leaving all the causal paths open.
- If some variables are not observed, things get more complicated, but do-calculus can tell you if the effect can be estimated from observational data alone.




Sufficient Conditioning Sets




Sufficient Conditioning Sets




Sufficient Conditioning Sets







Relation of interest




Confounder




Collider




Mediator




Independent




Cause of X




Consequence of X




Cause of Y



Consequence of Y




Unobserved confounder



$$X_1 \not\perp\!\!\!\perp X_4 \mid X_3$$

Unobserved confounder



$$X_1 \perp X_4 \mid X_3$$

Beyond NOTEARS

- **NOTEARS is linear:**

Find A

To minimize $\text{Mean}_i \|X_i - AX_i\|_F^2 + \lambda \|A\|_1$

Such that $\text{Trace } e^{A \odot A} = d$

- **It can be made non-linear:**

Find A

g_1, g_2 (pointwise)

To minimize $\text{Mean}_i \|X_i - g_2(Ag_1(X_i))\|_F^2 + \lambda \|A\|_1$

Such that $\text{Trace } e^{A \odot A} = d$